

Understanding Uncertainty Metrics

This report explains the uncertainty metrics and visualizations for 5 images analyzed with VAE-UNet.


Mean Expected Calibration Error (ECE): 0.0008

Mean Dice Score: 0.4938

Mean Brier Score: 0.0009

Mean Sparsification Error: 0.6701

Understanding Reliability Diagrams



RELIABILITY DIAGRAM EXPLANATION:

A reliability diagram shows how well a model's predicted probabilities match actual outcomes.

- Blue bars: The actual frequency of positive pixels in each confidence bin
- Green bars: The mean predicted probability (confidence) for each bin
- Gray histogram: Distribution of predictions across confidence levels
- Red lines: Highlight gaps between confidence and actual frequency
- Diagonal line: Perfect calibration (confidence = actual frequency)

INTERPRETATION:

- When blue bars are higher than green bars: Model is underconfident
- When green bars are higher than blue bars: Model is overconfident
- Expected Calibration Error (ECE): Weighted average of gaps between bars
 - Lower ECE values (closer to 0) indicate better calibration

ECE Values Interpretation:


- < 0.01: Excellent calibration
- 0.01-0.05: Good calibration
- 0.05-0.15: Fair calibration
- > 0.15: Poor calibration

WHY IT MATTERS:

Good calibration means the confidence values from your model are reliable. For medical segmentation, this helps clinicians know when to trust the model's predictions.

Understanding Sparsification Curves

Example Sparsification Curve



SPARSIFICATION CURVE EXPLANATION:

A sparsification curve shows whether your model's uncertainty estimates correlate with actual errors.

- Blue dashed line: Error when removing pixels randomly
- Red solid line: Error when removing pixels with highest uncertainty first
 - Green/Red fill: Area between curves (Sparsification Error)
- Red dot: Fraction of pixels that must be removed to halve the error

INTERPRETATION:

- If red line is below blue line (green area): Good uncertainty estimates!
This means removing high-uncertainty pixels reduces error faster than random removal.

- If red line is above blue line (red area): Poor uncertainty estimates.
Your model's uncertainty doesn't correlate well with actual errors.

- Sparsification Error (SE): Area between the curves
 - Positive SE: Good uncertainty estimates
 - Negative SE: Poor uncertainty estimates
- Larger positive values indicate better uncertainty quality

WHY IT MATTERS:

Good uncertainty estimates help identify which predictions might be wrong and where the model needs human verification in clinical applications.

Understanding the Visualization Plots

CORRELATION MATRIX EXPLANATION:

The correlation matrix shows how different metrics relate to each other:

- Values close to 1: Strong positive correlation (one increases, the other increases)
- Values close to -1: Strong negative correlation (one increases, the other decreases)
 - Values close to 0: Little or no correlation

Key relationships to look for:

- Dice Score vs. Uncertainty Metrics: Does better performance correlate with better calibration?
 - ECE vs. Sparsification Error: Do different uncertainty metrics agree with each other?

CALIBRATION ANALYSIS PLOT EXPLANATION:

This plot helps understand the pattern of calibration errors:

- X-axis: Maximum Calibration Error (MCE) - the largest calibration error in any bin
- Y-axis: Mean Absolute Calibration Error (MACE) - the average calibration error
- Color: Expected Calibration Error (ECE) - weighted average of calibration errors
 - Size: Dice Score - larger points indicate better segmentation performance

Interpretation by location:

- Points near the diagonal: Errors are consistent across all confidence levels
 - Points below diagonal: Errors concentrated in specific confidence bins
 - Bottom-left corner: Best calibration overall (low errors)
- Larger points in bottom-left: Ideal models (good performance, good calibration)

PAIRPLOT EXPLANATION:

The pairplot shows the relationships between all pairs of metrics:

- Diagonal: Distribution of each individual metric
- Off-diagonal: Scatter plots showing relationship between pairs of metrics

TEMPERATURE SCALING:

If your model has an ECE of 0.005, that's excellent calibration!

With temperature=2.0 giving better results, this suggests your model was slightly overconfident at the default temperature (T=1.0). Higher temperatures 'soften' predictions, making very confident predictions less extreme.